

Operating System Interview Questions:

Basic Questions:

1. **What is an Operating System?**
 - An OS is software that manages hardware and software resources on a computer and provides services for computer programs.
2. **What are the types of Operating Systems?**
 - **Batch OS, Time-Sharing OS, Distributed OS, Network OS, Real-Time OS, and Embedded OS.**
3. **What is the role of the kernel in an OS?**
 - The kernel manages system resources, including the CPU, memory, and hardware devices. It provides essential services like process management and memory management.
4. **What is a process in an OS?**
 - A process is a program in execution. It has its own memory space and system resources.
5. **What is the difference between a process and a thread?**
 - A **process** is an independent program with its own memory space, while a **thread** is a smaller unit of a process that shares memory and resources with other threads in the same process.
6. **What is multitasking?**
 - Multitasking is the ability of an OS to execute multiple tasks or processes concurrently.
7. **What is a deadlock?**
 - A deadlock occurs when two or more processes are unable to proceed because each is waiting for the other to release resources.
8. **What are the four necessary conditions for a deadlock?**
 - **Mutual Exclusion, Hold and Wait, No Preemption, and Circular Wait.**
9. **What is process scheduling?**
 - Process scheduling is the method by which the OS decides which process runs at a given time.
10. **What are the types of process scheduling algorithms?**

- **First-Come, First-Served (FCFS), Shortest Job Next (SJN), Round Robin (RR), Priority Scheduling, and Multilevel Queue Scheduling.**

11. What is virtual memory?

- Virtual memory allows programs to access more memory than physically available by using disk space as an extension of RAM.

12. What is a page fault?

- A page fault occurs when a program accesses a page that is not currently in memory.

13. What is paging in OS?

- Paging is a memory management scheme that eliminates the need for contiguous allocation of physical memory.

14. What is segmentation in OS?

- Segmentation divides memory into segments based on logical divisions like functions, arrays, etc.

15. What is thrashing?

- Thrashing occurs when the OS spends most of its time swapping data between memory and disk, leading to a significant performance slowdown.

Intermediate Questions:

16. What is a file system?

- A file system is a method of storing and organizing files on a storage device like a hard disk or SSD.

17. What are the types of file systems?

- **FAT, NTFS, ext4, HFS, APFS, and exFAT.**

18. What is the difference between FAT32 and NTFS?

- **FAT32** is older and has limitations like 4GB file size limit, while **NTFS** supports larger files, better security, and more features.

19. What is a system call?

- A system call is a request made by a program to the kernel for a service like file manipulation, process control, or communication.

20. What is the difference between a hard link and a soft link?

- A **hard link** points to the same inode, while a **soft link** (symlink) is a pointer to the file name.

21. What is a semaphore?

- A semaphore is a synchronization primitive used to control access to shared resources in a concurrent system.

22. What is the difference between binary and counting semaphores?

- A **binary semaphore** can only have values 0 or 1, while a **counting semaphore** can have a range of values.

23. What is a mutex?

- A mutex (mutual exclusion) is a locking mechanism used to ensure that only one thread can access a resource at a time.

24. What is a context switch?

- A context switch occurs when the OS switches from one process or thread to another, saving the state of the current process and loading the state of the next.

25. What is a shell in an OS?

- A shell is a user interface that allows interaction with the OS through commands, typically either command-line or graphical.

26. What is an inode in Linux?

- An inode is a data structure in the filesystem that stores information about a file, such as its size, permissions, and location on disk.

27. What is memory management?

- Memory management is the process of managing the computer's memory resources, including allocation, deallocation, and optimization of memory use.

28. What is the difference between a primary memory and a secondary memory?

- **Primary memory** (RAM) is fast and volatile, while **secondary memory** (HDD, SSD) is slower but non-volatile and used for long-term storage.

29. What is the difference between user mode and kernel mode?

- **User mode** is the mode in which applications run, while **kernel mode** is where the OS kernel executes and has full access to hardware.

30. What is a device driver?

- A device driver is software that allows the OS to communicate with hardware devices like printers, keyboards, and network interfaces.

Advanced Questions:

31. What is RAID?

- RAID (Redundant Array of Independent Disks) is a data storage technology that combines multiple disk drives for redundancy, performance, or both.

32. What is the difference between RAID 0, RAID 1, and RAID 5?

- **RAID 0:** Stripes data across disks for performance (no redundancy).
- **RAID 1:** Mirrors data across disks for redundancy.

- **RAID 5:** Uses striping and parity for redundancy and performance.

33. What is the difference between preemptive and non-preemptive scheduling?

- **Preemptive** scheduling allows the OS to interrupt a running process to assign CPU time to another process, while **non-preemptive** scheduling lets a process run to completion.

34. What is a socket in OS?

- A socket is an endpoint for communication between two machines over a network, commonly used in client-server applications.

35. What is a critical section?

- A critical section is a part of the code that accesses shared resources and must be executed by only one process or thread at a time.

36. What is an interrupt?

- An interrupt is a mechanism that temporarily halts the current execution of a program to allow the OS to respond to an event, like input from a user or hardware failure.

37. What is an event-driven system?

- An event-driven system is one where the flow of the program is determined by events like user actions, sensor outputs, or messages from other programs.

38. What is the purpose of a buffer?

- A buffer temporarily stores data to handle the differences in speed between two processes or devices (e.g., reading data from disk to memory).

39. What is a bootloader?

- A bootloader is a small program that loads the OS into memory when a computer starts up.

40. What is a system call interface (SCI)?

- The SCI is the interface between user applications and the kernel, enabling applications to request OS services.

41. What is the difference between a process and a program?

- A **process** is an executing instance of a program, while a **program** is a static set of instructions stored on disk.

42. What is memory-mapped I/O?

- Memory-mapped I/O is a method of performing input/output operations by mapping I/O devices directly into the system's memory space.

43. What is the difference between a signal and a semaphore?

- A **signal** is a notification mechanism for inter-process communication, while a **semaphore** is used for synchronizing access to shared resources.

44. What is the role of a file descriptor?

- A file descriptor is a unique identifier used by the OS to access and manage open files or I/O streams.

45. What is load balancing?

- Load balancing is the process of distributing tasks across multiple resources (e.g., servers) to ensure optimal resource utilization and prevent overloading.

46. What is a race condition?

- A race condition occurs when multiple processes or threads attempt to access shared resources simultaneously, leading to unpredictable results.

47. What is the purpose of a cache in memory management?

- A cache stores frequently accessed data to speed up access and reduce the load on slower storage devices.

48. What is process synchronization?

- Process synchronization ensures that processes or threads coordinate their actions to prevent conflicts when accessing shared resources.

49. What is the purpose of the fork() system call?

- The fork() system call creates a new process by duplicating the calling process.

50. What is a page replacement algorithm?

- A page replacement algorithm determines which pages to swap in and out of memory when there is a page fault and memory is full (e.g., **LRU**, **FIFO**, **Optimal**).